

Setting up a Data Warehouse using Amazon Redshift Serverless

Use case

Any Company Events is ready for you to create an Amazon Redshift Serverless data warehouse. After you create a namespace and workgroup, you will create a schema and tables that will be loaded with TICKET data. Finally, you will create a view and query a sample of the data.'

In this lab, you create an Amazon Redshift Serverless namespace and workgroup, schema, and tables. You load the tables with sample data and then query the data.

Objectives

By the end of this lab, you should be able to do the following:

- Create a data warehouse with Amazon Redshift Serverless.
- Create a schema.
- Create a table.
- Load the table with sample data.

On the Get started with Amazon Redshift Serverless page, configure the following.

In the Configuration section, select Customize settings.

In the Namespace section, configure:

- For Target namespace, enter `lab-namespace`.
- For Admin user credentials, select Customize admin user credentials.
 - For Admin user name, enter `dbadmin`.
 - For Admin password, select Manage admin credentials in AWS Secrets Manager.

In the Associated IAM roles section, configure:

- Choose **Associate IAM role**.
- Select the IAM Role that contains the string RedshiftAccessRole in name.
- Choose **Associate IAM roles**.
- Select RedshiftAccessRole to make default IAM role.
- Choose **Set default**.
- Select Make default.
- In the Make default IAM role prompt, choose **Confirm**.

In the Workgroup section, configure:

- For Workgroup name, enter `lab-workgroup`.
- For Virtual private cloud (VPC), select `LabVPC`.
- For VPC security groups, select the security group with Enable access to redshift in description.
- For Subnet, select Private Subnet 1, Private Subnet 2, Private Subnet 3.

Choose **Save configuration**.

A new window is displayed with a message *It may take a few minutes to complete. After completing the setup, you can work with your data.*

Note: Wait for the status to display **Completed**.

Choose **Continue**.

Get started with Amazon Redshift Serverless [Info](#)

To start using Amazon Redshift Serverless, set up your serverless data warehouse and create a database. When you create and use your serverless data warehouse for the first time, a \$300 credit toward Redshift Serverless usage is applied to the account.

Configuration



Use default settings

Default settings have been defined to help you get started. You can change them at any time later.



Customize settings

Customize your settings for your specific needs.

Namespace [Info](#)

Namespace is a collection of database objects and users. Data properties include database name and password, permissions, and encryption and security.

Target namespace

This is a unique name that defines the namespace.

The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).

▼ Database name and password

Database name

The name of the first database in the Amazon Redshift Serverless environment.

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word.

Admin user credentials

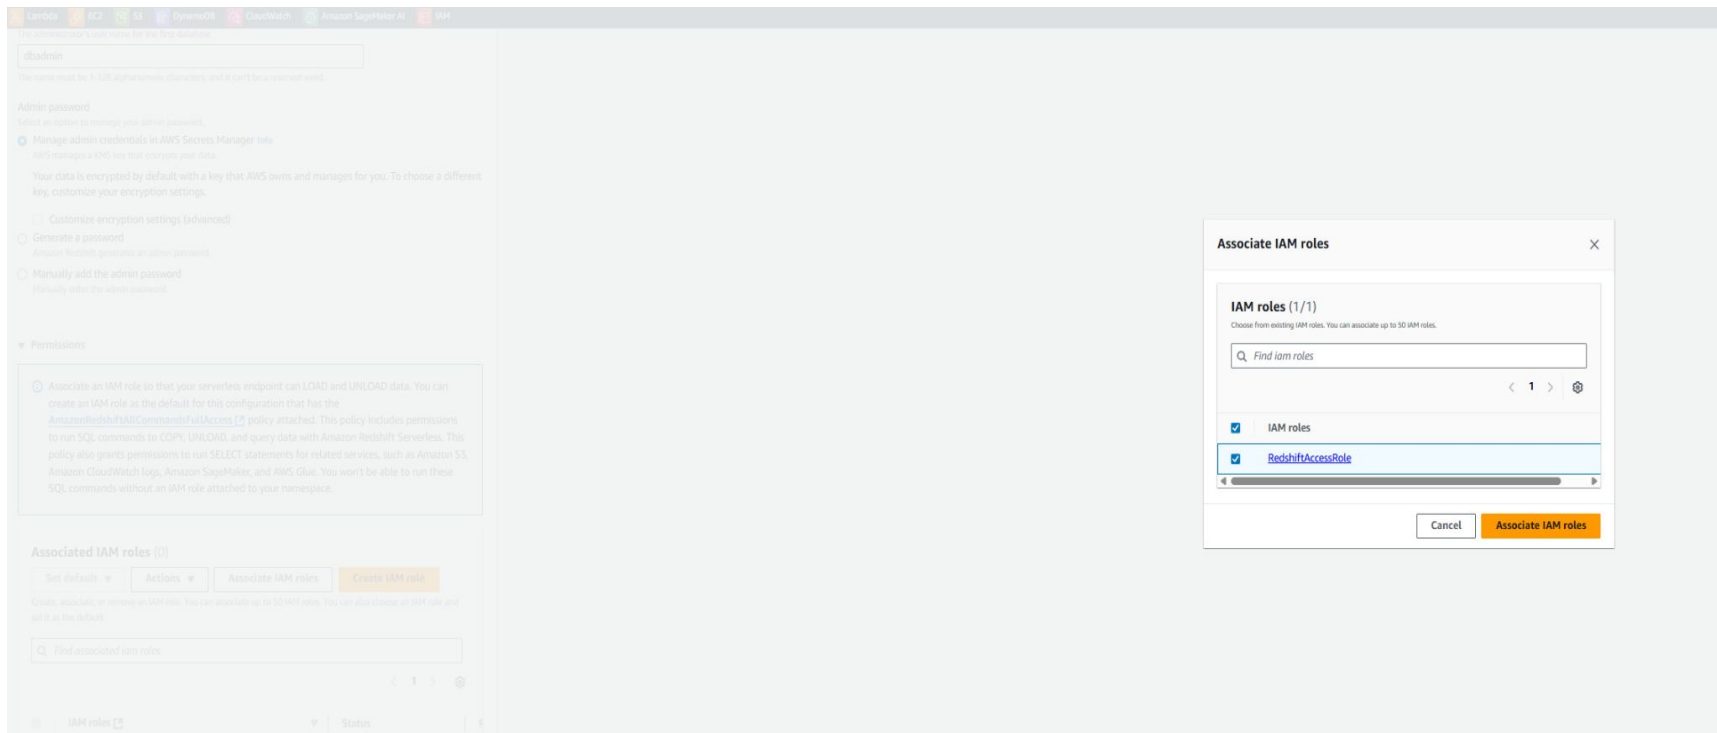
IAM credentials provided as your default admin user credentials. To add a new admin username and password, customize admin user credentials.

☒ Customize admin user credentials

To use the default IAM credentials, clear this option.

Admin user name

The administrator's user name for the first database.



Create a schema

In this task, you create a schema in the current database.

A schema is a namespace that contains named database objects such as tables, views, and user-defined functions (UDFs). A database can contain one or multiple schemas, and each schema belongs to only one database. Two schemas can have different objects that share the same name.

Note: Amazon Redshift automatically creates a schema called public for every new database. When you don't specify the schema name while creating database objects, the objects go into the public schema.

Connect to workgroup

13. On the Serverless dashboard page, choose **Query data**.
The query editor now opens in a new browser tab.
14. In the Redshift query editor v2 pane, choose the vertical ellipses (three dots) at Serverless:lab-workgroup.
15. Choose **Create connection**.
16. On the Connect to lab-workgroup window prompt, choose **AWS Secrets Manager**.
17. Select **Choose a secret** and choose redshift!!lab-namespace-dbadmin.
18. Choose **Create connection**.
19. Command: To create a new schema named US_SALES, run the following query:
The following command either creates the *US_SALES* schema for the dev database, or does nothing and returns a message if it already exists.
This sql command defines a new schema for the *dev* database.

Create schema if not exists us_sales;

To create a table in *us_sales* schema, run the following command:

```
CREATE TABLE us_sales.supplier (  
SupplierID int,  
City varchar (255)  
);
```

To Insert

```
INSERT INTO us_sales.supplier VALUES (781, 'San Jose'), (990, 'Palo Alto');
```

To view the data inserted into a table, run the following command:

```
SELECT * from us_sales.supplier;
```

Load sample data from Amazon S3

In this task, you use the COPY command to load data into the table.

The COPY command uses the Amazon Redshift massively parallel processing (MPP) architecture to read and load data in parallel from files in Amazon S3, from an Amazon DynamoDB table, or from text output from one or more remote hosts.

Note: The COPY command appends the new input data to any existing rows in the table. The maximum size of a single input row from any source is 4 MB.

Copy edit: To load data into the table from the Amazon S3 bucket, replace the S3_BUCKET_PATH placeholder value with the S3BucketPath value listed to the left of these instructions. Then, run the following query:

```
COPY us_sales.supplier
```

```
FROM 'S3_BUCKET_PATH/supplier.csv'
```

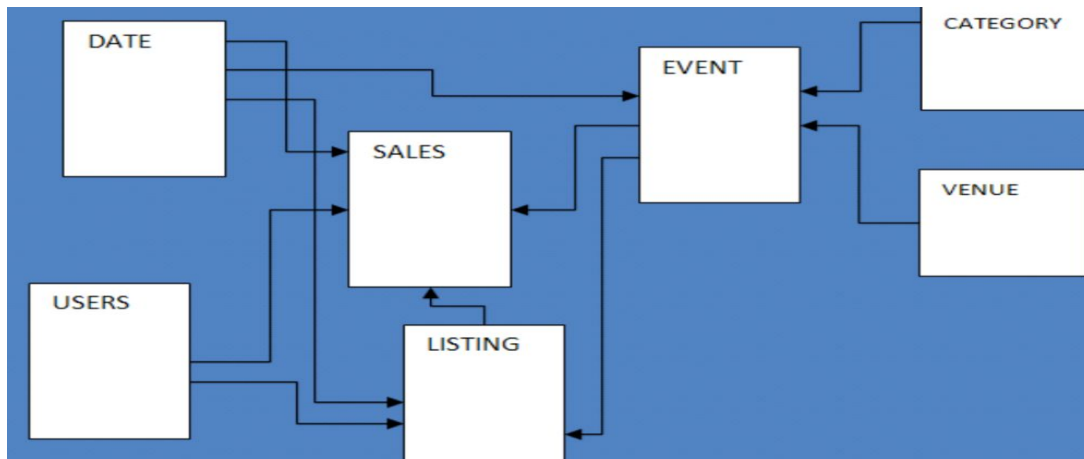
```
24.   iam_role default csv;
```

```
SELECT * FROM us_sales.supplier;
```

Ingest and query TICKIT data

AnyCompany Events wants to build database application that helps analysts track sales activity from TICKIT website where users buy and sell tickets online for sporting events, shows, and concerts.

In this task, you load the TICKIT dataset from Amazon S3 to Redshift Serverless database and build queries to help analysts to use this information to provide incentives to buyers and sellers who frequent the site, to attract new users, and to drive advertising and promotions.



To create tables in the dev database, individually copy and run the following create table statements:

```
CREATE TABLE us_sales.users (userid INTEGER NOT NULL DISTKEY SORTKEY, username CHAR(8), firstname VARCHAR(30), lastname VARCHAR(30), city VARCHAR(30), state CHAR(2), email VARCHAR(100), phone CHAR(14), likesports BOOLEAN, liketheatre BOOLEAN, likeconcerts BOOLEAN, likejazz BOOLEAN, likeclassical BOOLEAN, likeopera BOOLEAN, likerock BOOLEAN, likevegas BOOLEAN, likebroadway BOOLEAN, likemusicals BOOLEAN);
```


Create all the other tables

```
CREATE TABLE us_sales.venue (venueid SMALLINT NOT NULL DISTKEY SORTKEY, venuename VARCHAR(100), venuecity VARCHAR(30), venuestate CHAR(2), venuesseats INTEGER);
```

```
CREATE TABLE us_sales.category (catid SMALLINT NOT NULL DISTKEY SORTKEY, catgroup VARCHAR(10), catname VARCHAR(10), catdesc VARCHAR(50));
```

```
CREATE TABLE us_sales.date (dateid SMALLINT NOT NULL DISTKEY SORTKEY, caldate DATE NOT NULL, day CHAR(3) NOT NULL, week SMALLINT NOT NULL, month CHAR(5) NOT NULL, qtr CHAR(5) NOT NULL, year SMALLINT NOT NULL, holiday BOOLEAN DEFAULT('N'));
```

```
CREATE TABLE us_sales.event (eventid INTEGER NOT NULL DISTKEY, venueid SMALLINT NOT NULL, catid SMALLINT NOT NULL, dateid SMALLINT NOT NULL SORTKEY, eventname VARCHAR(200), starttime TIMESTAMP);
```

```
CREATE TABLE us_sales.event (eventid INTEGER NOT NULL DISTKEY, venueid SMALLINT NOT NULL, catid SMALLINT NOT NULL, dateid SMALLINT NOT NULL SORTKEY, eventname VARCHAR(200), starttime TIMESTAMP);
```

```
CREATE TABLE us_sales.listing (listid INTEGER NOT NULL DISTKEY, sellerid INTEGER NOT NULL, eventid INTEGER NOT NULL, dateid SMALLINT NOT NULL SORTKEY, numtickets SMALLINT NOT NULL, priceperticket DECIMAL(8,2), totalprice DECIMAL(8,2), listtime TIMESTAMP);
```

```
CREATE TABLE us_sales.sales (salesid INTEGER NOT NULL, listid INTEGER NOT NULL DISTKEY, sellerid INTEGER NOT NULL, buyerid INTEGER NOT NULL, eventid INTEGER NOT NULL, dateid SMALLINT NOT NULL SORTKEY, qty sold SMALLINT NOT NULL, pricepaid DECIMAL(8,2), commission DECIMAL(8,2), saletime TIMESTAMP);
```

To load TICKIT dataset from the Amazon S3 bucket, replace the S3_BUCKET_PATH placeholder value with the S3BucketPath created ex s3://labstack-502883f7-b4ce-49fe-bc70-c4f0be-databucket-afjnmqcyqoxv/tickitdb. Then, run the following query:

```
copy us_sales.users from 'S3_BUCKET_PATH/allusers.csv' iam_role default csv;
```

```
copy us_sales.venue from 'S3_BUCKET_PATH/venue_pipe.txt' iam_role default delimiter '|';
```

```
copy us_sales.category from 'S3_BUCKET_PATH/category_pipe.txt' iam_role default delimiter '|';
```

```
copy us_sales.date from 'S3_BUCKET_PATH/date2008.csv' iam_role default csv;
```

```
copy us_sales.event from 'S3_BUCKET_PATH/allevnts_pipe.txt' iam_role default delimiter '|' timeformat 'YYYY-MM-DD HH:MI:SS' ;
```

```
copy us_sales.listing from 'S3_BUCKET_PATH/listings_pipe.txt' iam_role default delimiter '|';
```

```
copy us_sales.sales from 'S3_BUCKET_PATH/sales_tab.txt' iam_role default delimiter '\t' timeformat 'MM/DD/YYYY HH:MI:SS' ,
```